

Research Statement

CIFRE PhD — CREST - IMT, Generali France

*Regime-Conditional Generative Models, Volterra Capital Proxies
and Mean Field Game Lapse Dynamics for Insurance Asset Allocation*

Candidate :	Yassine Mannai
Contact :	yassine.mannai@dauphine.eu · +33 (0)6 23 64 31 80
Programme :	MSc in Quantitative Finance (272), Paris Dauphine-PSL
Academic supervisors :	Caroline Hillairet (CREST, ENSAE Paris) and Anthony Réveillac (INSA Toulouse, IMT)
Industry supervisor :	Mamy Ramamonjisoa — SAA team, Generali France
Investment Division :	Cedrik d'Aviau de Ternay — Chief Investment Officer, Generali France
Start date :	January 2027 · Duration : 36 months

Thesis in one sentence.

This thesis introduces the first regime-conditional generative framework with certified W_2 guarantees under online regime inference, coupled to a differentiable functional capital proxy, to enable a fully dynamic, capital-aware Total Portfolio Approach for a life insurer under Solvency II. The three axes form a single closed control loop : a certified scenario generator (Axis 1) feeds a fast differentiable SCR proxy (Axis 2), whose gradient constraint drives the MFG-based lapse modelling and capital allocation programme (Axis 3).

1. INDUSTRIAL CONTEXT AND MOTIVATION

1.1. The 2022 dependence crisis and the limits of static allocation

The year 2022 was a structural stress test for insurance asset managers : unlike earlier single-asset shocks, equities and bonds fell simultaneously and diversification broke down. For a life insurer such as Generali France, this exposed a structural flaw in Strategic Asset Allocation (SAA), which assumes stable, linear dependence between economic factors. The equity/bond correlation switched from -0.3 (diversifying) to $+0.6$ (concentration risk) within weeks — a regime reversal that fixed-correlation models cannot anticipate.

This motivates the central problem of this thesis : *how to characterise and optimally control the solvency position of an insurer when the joint dynamics of economic factors are governed by a latent Markov chain inducing non-stationary, non-linear dependence structures across regimes ?*

1.2. Three compounded bottlenecks in the current SAA process

The SAA process at Generali France runs in three sequential steps, each with structural limitations. **(1) Economic Scenario Generator (ESG)** — a proprietary ESG produces thousands of trajectories (rates, spreads, equity, real estate, inflation) calibrated with **fixed correlations averaged across regimes**, feeding all downstream actuarial calculations. **(2) SCR via nested LSMC** — the Solvency Capital Requirement, the 99.5% one-year VaR of net asset value change, requires the Best Estimate of Liabilities (BEL) under 10,000+ scenarios ; Least-Squares Monte Carlo (LSMC, Krah et al., 2018) replaces the inner simulation by an offline polynomial proxy. **(3) Annual Markowitz optimisation** — the SCR is a fixed constraint in a mean-variance optimisation revised **once per year**, with no automatic identification of the dominant portfolio objective (duration, inflation, carry, capital efficiency), even though the relevant factor varies across portfolios and regimes.

Four structural limitations.

1. **Regime-blind dependence.** Fixed correlations cannot anticipate regime reversals : a Markowitz optimum for a low-inflation regime becomes high-risk under stagflation.
2. **LSMC proxy limitations.** The polynomial proxy is accurate only within its regime-agnostic calibration region, degrading precisely in the stress zones most relevant for capital decisions, and is not expressed in common economic factors.
3. **Static, siloed decision-making.** Allocation and hedging are decided separately — suboptimal for an insurer with convex critical losses — and the objective is imposed uniformly rather than inferred from each portfolio’s dominant risk factor.
4. **Private asset valuation bias.** Private equity, private debt and real estate NAVs are smoothed quarterly : in 2008 private equity NAVs stayed stable for 6–12 months before correcting sharply, understating stress-regime dependence and introducing a lag bias in the SCR.

1.3. The Total Portfolio Approach : target framework

The **Total Portfolio Approach (TPA)** evaluates each decision through its marginal contribution to total balance-sheet risk via common factors (Duration, Inflation, Growth, Liquidity, Carry). Implementations at sovereign wealth funds (Future Fund, CPP Investments) report about +1% excess return per year over classical SAA (Thinking Ahead Institute, 2025 ; CAIA, 2024). For an insurer under Solvency II, the factor space must also include liability-side factors (surrender, longevity, yield-curve/BEL interactions), and every decision must internalise a capital cost. Existing actuarial tools address none of this jointly ; this thesis fills the gap along three coupled axes : regime-conditional factor simulation (Axis 1), fast capital proxy (Axis 2), and MFG-based lapse modelling and capital allocation (Axis 3).

2. RESEARCH OVERVIEW




FIGURE 1 – The three axes form a single closed capital-aware control loop : certified regime-conditional scenarios feed a fast differentiable capital proxy, whose gradient drives the capital allocation programme ; realised regimes feed back into the generator.

Approach	Regime-aware	Differentiable	W_2/L^2 certified	Insurance-calibrated
Static SAA / Markowitz	✗	✗	✗	✓
Copulas (Sklar, Fermanian)	✗	✗	✓	✗
LSMC proxy (Krah et al.)	✗	✗	✓	✓
GAN/VAE time-series (Wiese, Yoon)	✗	✓	✗	✗
Score-based models (Ocello et al.)	✗	✓	✓	✗
Classical lapse models (actuarial)	✗	✗	✗	✓
This thesis (Axes 1–3)	✓	✓	✓	✓

TABLE 1 – Positioning of the research programme against the state of the art.

✓ = requirement met, ✗ = not addressed.

3. RESEARCH PROGRAMME

3.1. Axis 1 — Regime-Conditional Factor Dependence

Research question.

Research Question 1. *How to generate the joint distribution of TPA factors conditionally on a sequentially estimated macroeconomic regime, while providing : (i) formal goodness-of-fit guarantees on the learned dependence structure ; and (ii) non-asymptotic W_2 -convergence bounds under particle-filter inference error on the regime state ?*

State of the art.

Regime identification. Hidden Markov Models (Hamilton, 1989) and Sequential Monte Carlo (SMC) particle filters (Chopin & Papaspiliopoulos, 2020) provide tractable real-time estimation of latent macroeconomic regimes $z_t \in \{1, \dots, K\}$. Their combination — an online HMM-SMC filter — yields a posterior $p(z_t | \mathcal{F}_t)$ updated at each observation.

Generative Adversarial Networks. Applied to financial series by Wiese et al. (2020, Quant-GAN) and Yoon et al. (2019, TimeGAN), conditional GANs generate $\mathbf{F}_{t+1} | z_t = k \sim G_\theta(\varepsilon, k)$ with full flexibility (no parametric dependence assumption), but suffer from *mode collapse* and the absence of distributional guarantees, limiting regulatory auditability.

Variational Autoencoders. CVAEs encode each trajectory into a regime-conditional posterior and maximise the ELBO ; they offer stable training and an auditable latent structure, but produce overly smooth samples due to the Gaussian posterior approximation.

Score-Based Generative Models (SGMs). Strasman, Ocello et al. (TMLR, 2024) establish W_2 convergence bounds :

$$\text{KL}(\hat{p} \| p^*) \leq \int_0^T g(t)^2 \varepsilon_{\text{score}}(t) dt + \varepsilon_{\text{init}}. \quad (1)$$

Gentiloni-Silveri & Ocello (ICML, 2025) extend these guarantees to non-log-concave distributions (Gaussian mixtures), directly applicable to multi-regime market models. De Marco, Pham & Zanni (2025) further introduce jump-diffusion Schrödinger bridges to capture discontinuous regime transitions.

Copulas : a validation framework, not a generative model. Sklar’s theorem separates margins from dependence, and Fermanian (2013, 2022) provides goodness-of-fit tests for conditional copulas. Classical parametric families (Gaussian, Student- t , Archimedean), however, impose rigid tails, scale poorly in high dimension and cannot capture cross-regime dynamics. Copulas therefore serve here exclusively as a *validation tool* : the conditional tail-dependence coefficient λ_k is estimated and tested via Fermanian’s specification tests to validate generative outputs, not to generate them.

Identified gaps. Neither conditional GANs, VAEs nor SGMs have been applied to the joint distribution of TPA factors with real-time regime identification. For the GAN/VAE : the absence of formal regime-conditional tail-dependence validation. For the SGM : the absence of W_2 bounds for the conditional case with sequentially estimated regime states. These constitute the theoretical contribution of Axis 1.

Proposed approach.

Stage 1 (Year 1) — cGAN and cVAE. Let z_t be identified by an online HMM-SMC filter. A conditional GAN $G_\theta(\varepsilon, k)$ and a CVAE are trained in parallel, both incorporating the private asset vintage lag factor $\tilde{F}_t^{\text{PA}} = \text{NAV}_{t-1}^{\text{PA}}$ to capture the quarterly smoothing bias. Model comparison is conducted on : Wasserstein-2 distance W_2 , Maximum Mean Discrepancy (MMD), per-regime tail calibration via Fermanian’s λ_k tests, autocorrelation structure, and regulatory auditability.

Stage 2 (Year 2) — Conditional SGM with W_2 guarantees. A conditional OU forward process for regime k :

$$d\mathbf{x}^\tau = -\mathbf{x}^\tau d\tau + \sqrt{2} g_k(\tau) dW_\tau, \quad \tau \in [0, T], \quad (2)$$

with regime-specific noise schedule $g_k(\tau)$. The theoretical contribution is the extension of the W_2 bounds of Gentiloni-Silveri & Ocello (ICML 2025) to the conditional case with online SMC regime estimation —

quantifying how inference error in $p(z_t|\mathcal{F}_t)$ propagates into the generative bound. A scenario selection layer then extracts a central regime-conditional scenario for downstream reporting and optimisation diagnostics.

3.2. Axis 2 — Supervised Capital Proxy

Research question.

Research Question 2. *How to construct a proxy for the insurance SCR — expressed as a function of TPA factors, the current regime, and liability states including private asset NAV lags formalised as a Volterra functional — that is accurate in decision-relevant zones with controlled L^2 error bounds per regime and sub-10ms inference time, and exposes a differentiable interface for its direct embedding as an explicit capital constraint in a downstream stochastic optimisation programme ?*

State of the art.

Classical proxy models. The LSMC framework (Krah et al., 2018) approximates the BEL by polynomial regression on a design-of-experiment. Its two critical limitations are regime-blindness and degraded accuracy in stress zones. Neural network proxies (Huber & Schmocker, 2022) achieve 1–3% accuracy at $< 10\text{ms}$ inference. Bourgey (2020) provides theoretical foundations via multi-level Monte Carlo (MLMC) estimators for nested expectations and polynomial chaos expansions with L^2 error estimates, motivating the adaptive-weighting architecture of Axis 2.

Recent generative and flow-based surrogates. Three recent architectures merit systematic comparison for their uncertainty-quantification and stability properties : **Neural ODEs** (Chen et al., 2018), which model the proxy surface as a continuous-depth dynamical system ; **Conditional Flow Matching** (Lipman et al., 2022), which learns straight-line ODE paths offering stable training and a full conditional density $p(\widehat{\text{SCR}}|\mathbf{F}, z, \mathbf{P})$ for uncertainty-aware allocation ; and **score-based diffusion surrogates**. **Tabular diffusion models** (TabDDPM, Kotelnikov et al., 2023) additionally serve as design-of-experiment augmentation for under-represented regimes.

Volterra processes and path-dependent risk measures. The private asset NAV lag is not a Markovian state variable. The observed smoothed value $\widehat{\text{NAV}}_t = \int_0^t K(t, s) \text{NAV}_s^* ds$ is a Volterra functional of the true (unobservable) NAV process, where $K(t, s)$ is a smoothing kernel. The theory of affine Volterra processes (Abi Jaber, Larsson & Pulido, 2019 ; El Euch & Rosenbaum, 2019) provides a tractable framework for such path-dependent dynamics. This motivates extending the capital risk measure to the functional setting :

$$\rho^V(L_T) = \phi \left(\int_0^T K(T, s) \text{CVaR}_\alpha(L_s) ds \right), \quad (3)$$

where ϕ is a monetary risk measure. Equation (3) formalises the temporally diffused risk crystallisation of private assets and provides a theoretical grounding for the vintage lag correction embedded in the proxy state vector.

Model comparison and interoperability metrics. Proxy architectures will be evaluated on two sets of criteria. *Comparison metrics* : (i) accuracy — RMSE, relative error at the SCR regulatory threshold, L^∞ norm in stress zones ; (ii) calibration — prediction-interval coverage, reliability diagrams ; (iii) speed — inference latency (target $< 5\text{ms}$) and batch throughput ; (iv) interpretability — SHAP values, integrated gradients, first-order Sobol sensitivity indices $S_i = \text{Var}[\mathbb{E}[\widehat{\text{SCR}}|\mathbf{F}_i]]/\text{Var}[\widehat{\text{SCR}}]$. *Interoperability metrics* : (i) differentiability — $\nabla_{\mathbf{F}}\widehat{\text{SCR}}$ available via automatic differentiation for gradient-based embedding in the Axis 3 optimiser ; (ii) API compatibility with GPS and PyTorch/JAX optimisation frameworks ; (iii) monotonicity and convexity preservation as required by Solvency II constraints ; (iv) memory footprint for production deployment.

Identified gaps. No existing proxy is jointly conditioned on the macroeconomic regime, TPA factors and private asset NAV lag. None accounts for the path-dependent Volterra structure of NAV smoothing, and none provides the differentiable interface required for its direct embedding as an explicit capital constraint in a dynamic optimisation programme.

Proposed approach. We learn the supervised proxy :

$$\widehat{\text{SCR}}(t) = f_{\theta}(\mathbf{F}_t, z_t, \mathbf{P}_t), \quad (4)$$

where $\mathbf{P}_t$ includes effective duration, average guaranteed rate, observed lapse ratio and the vintage lag $\tilde{F}_{t-1}^{\text{PA}}$. The network f_{θ} is trained with **adaptive sample weighting** concentrating mass near the regulatory threshold and in stress regimes. The generalisation bound :

$$|\text{SCR}(t) - f_{\theta}(\mathbf{F}_t, z_t, \mathbf{P}_t)| \leq \varepsilon_k(n_k, \delta), \quad (5)$$

decreasing with n_k (training points in regime k) and with confidence $1 - \delta$. The four architectures (MLP+attention, Neural ODE, Flow Matching, score-based) are compared on the full metric set above, and the selected proxy is exported as a differentiable module ($\nabla_{\mathbf{F}} \widehat{\text{SCR}}$ available) for direct use as an explicit constraint in the Axis 3 programme.

3.3. Axis 3 — Mean Field Game Network Model for Lapse Risk and Capital Allocation

Research question.

Research Question 3. *How to model policyholder surrender dynamics as a mean field game on a heterogeneous policy network conditioned on the macroeconomic regime of Axis 1, derive the regime-conditional aggregate lapse process and its path-dependent impact on the SCR proxy of Axis 2, and solve the resulting capital allocation problem under a Solvency II shortfall constraint ?*

State of the art.

Classical lapse models. Current actuarial practice models surrender rates as deterministic functions of the interest-rate spread, calibrated on historical cohorts (Kochanski, 2010 style). Dynamic models incorporate rate-sensitivity but assume independence across policyholders — ignoring social contagion. In 2022, French life insurers observed correlated surrender waves far in excess of individual rational behaviour, motivating a network-contagion approach.

Mean field games. Lasry & Lions (2007) and Carmona & Delarue (2018) establish MFG theory for large populations of interacting agents : each agent solves an HJB equation whose coefficients depend on the aggregate distribution m_t (mean field), yielding a Nash equilibrium characterised by a coupled forward-backward PDE system. MFGs have been applied to systemic risk in banking (Carmona et al., 2015) but not to insurance lapse contagion.

Financial network contagion. Bardoscia et al. (2021) establish that network topology — degree distribution, community structure, contagion thresholds — determines systemic amplification. Applied to a policyholder network (nodes = contracts, edges = social or geographic proximity), the same framework governs surrender cascades.

Stochastic control and capital allocation. Réveillac (2013) and Pham (2009) provide the BSDE and HJB framework for optimal allocation under dynamic risk constraints. Rosazza Gianin et al. (2024) supply axiomatic capital allocation rules under shortfall risk measures, directly applicable once the MFG lapse process is embedded as a liability-side state variable.

Identified gaps. No existing model treats policyholder lapse as a MFG equilibrium on a network conditioned on the macroeconomic regime. No framework connects the aggregate lapse process to a differentiable SCR proxy or to a Solvency II capital allocation programme.

Proposed approach. Each policyholder i is a MFG agent with in-force indicator $x_i(t)$ and value function V^i satisfying

$$-\partial_t V^i = \min_{\alpha_i} \left[\ell(x_i, \alpha_i, m_t, z_t) + \mathcal{L}^{\alpha_i} V^i \right], \quad m_t = \frac{1}{N} \sum_j \delta_{x_j(t)}, \quad (6)$$

where z_t is the regime from Axis 1 and local interaction terms encode the network topology. The MFG equilibrium yields a regime-conditional aggregate lapse rate $\lambda^*(t, z_t, m_t)$ that enters the liability state vector $\mathbf{P}_t$ of the Axis 2 proxy. The capital allocation problem then reads

$$\min_{\mathbf{w}} \text{CVaR}_{0.99}(\Delta \text{NAV}_T) \quad \text{s.t.} \quad \widehat{\text{SCR}}(\mathbf{F}_t, z_t, \mathbf{P}_t(\lambda^*)) \leq \bar{C}, \quad (7)$$

where the differentiable proxy of Axis 2 enables gradient-based resolution of (7). The MFG system is solved via deep Galerkin methods (physics-informed NNs) and validated on the 2022 French surrender-wave episode.

4. TIMELINE (36 MONTHS)

- **Year 1 — Axis 1.** Bloomberg factor database with NAV lags; online HMM-SMC calibration; regime-conditional GAN/CVAE validated on 2008/2020/2022 crises. Publication in mathematical finance / statistics.
- **Year 2 — Axis 2.** Actuarial design-of-experiment (Prophet/MoSes) with TabDDPM augmentation; SCR proxy architecture comparison and Volterra L^2 bounds; differentiable proxy embedded as optimisation constraint. Publication in actuarial science / insurance mathematics.
- **Year 3 — Axis 3.** MFG-network lapse calibration on the 2022 French surrender-wave episode; full three-axis coupling; capital allocation out-of-sample backtesting; defence. Publication in quantitative finance.

5. CONCLUSION

This thesis delivers a tightly coupled three-module system designed to make the **Total Portfolio Approach operational for a life insurer under Solvency II**, with direct integration into Generali France’s proprietary platform GPS :

- **Module 1 (Axis 1)** replaces static ESG correlations with an auditable regime-conditional generative model, enabling real-time scenario generation that reflects the actual dependence structure of the prevailing macroeconomic regime, including the path-dependent correction dynamics of private assets.
- **Module 2 (Axis 2)** replaces the multi-hour SCR calculation with a $< 5\text{ms}$ differentiable proxy grounded in Volterra functional theory, making dynamic portfolio optimisation loops computationally feasible and the capital constraint analytically tractable as an explicit gradient-available term.
- **Module 3 (Axis 3)** replaces deterministic lapse tables with a mean field game equilibrium on the policyholder network, delivering a regime-conditional aggregate lapse process that feeds directly into the SCR proxy and enables a capital allocation programme with explicit Solvency II guarantees.

Beyond the industrial deliverables, the thesis makes three mathematically self-contained contributions : the extension of W_2 convergence theory to regime-conditional score-based models under sequential estimation error ; the formalisation of private asset valuation bias through the theory of Volterra processes and path-dependent risk measures ; and the characterisation of policyholder surrender equilibria as mean field games on network-structured populations with regime-conditional interaction kernels, and their integration into Solvency II capital allocation programmes via the differentiable SCR proxy of Axis 2. Together, these contributions establish rigorous foundations for a new generation of quantitative tools in insurance balance sheet management — tools that are simultaneously *theoretically grounded, computationally efficient, interoperable and regulatory-compliant*.

REFERENCES

- [1] Abi Jaber, E., Larsson, M., & Pulido, S. (2019). Affine Volterra processes. *The Annals of Applied Probability*, 29(6), 3155–3200.
- [2] Bardoscia, M., Barucca, P., Battiston, S., Caccioli, F., Cimini, G., Garlaschelli, D., Saracco, F., Squartini, T., & Caldarelli, G. (2021). The physics of financial networks. *Nature Reviews Physics*, 3(7), 490–507.

- [3] Bourgey, F. (2020). *Stochastic approximations for financial risk computations*. PhD thesis.
- [4] CAIA Association. (2024). *Innovation Unleashed : The Rise of Total Portfolio Approach*.
- [5] Carmona, R., & Delarue, F. (2018). *Probabilistic Theory of Mean Field Games with Applications*. Springer.
- [6] Chen, R.T.Q., Rubanova, Y., Bettencourt, J., & Duvenaud, D. (2018). Neural ordinary differential equations. *Advances in Neural Information Processing Systems*, 31.
- [7] Chopin, N., & Papaspiliopoulos, O. (2020). *An Introduction to Sequential Monte Carlo*. Springer.
- [8] De Marco, S., Pham, H., & Zanni, D. (2025). Schrödinger bridges with jumps for time series generation. *arXiv preprint arXiv :2602.20011*.
- [9] El Euch, O., & Rosenbaum, M. (2019). The characteristic function of rough Heston models. *Mathematical Finance*, 29(1), 3–38.
- [10] Fermanian, J.-D. (2013). An overview of the goodness-of-fit test problem for copulas. In *Copulae in Mathematical and Quantitative Finance*, Springer, 61–89.
- [11] Fermanian, J.-D. (2022). Conditional copulas, dependence modeling and the role of conditioning. *Annals of Statistics*.
- [12] Gentiloni-Silveri, M., & Ocello, A. (2025). Beyond log-concavity and score regularity : Improved convergence bounds for score-based generative models in W_2 -distance. *Forty-Second International Conference on Machine Learning (ICML)*.
- [13] Hamilton, J.D. (1989). A new approach to the economic analysis of nonstationary time series. *Econometrica*, 57(2), 357–384.
- [14] Huber, M., & Schmock, B. (2022). A neural network approach to efficient valuation of large portfolios of variable annuities. *Insurance : Mathematics and Economics*, 105, 96–116.
- [15] Kotelnikov, A., Baranchuk, D., Rubachev, I., & Babenko, A. (2023). TabDDPM : Modelling Tabular Data with Diffusion Models. *International Conference on Machine Learning (ICML)*.
- [16] Krah, A.-S., Nikolic, Z., & Korn, R. (2018). A Least-Squares Monte Carlo framework in proxy modeling of life insurance companies. *Risks*, 6(2), 62.
- [17] Lasry, J.-M., & Lions, P.-L. (2007). Mean field games. *Japanese Journal of Mathematics*, 2(1), 229–260.
- [18] Lipman, Y., Chen, R.T.Q., Ben-Hamu, H., Nickel, M., & Le, M. (2022). Flow Matching for Generative Modeling. *International Conference on Learning Representations (ICLR)*.
- [19] Pham, H. (2009). *Continuous-Time Stochastic Control and Optimization with Financial Applications*. Springer.
- [20] Réveillac, A. (2013). CRRA utility maximization under dynamic risk constraints. *Mathematical Finance*.
- [21] Rosazza Gianin, E., et al. (2024). Shortfall systemic risk measures, capital allocation rules and generalised collapse to the mean : Theory and practice. *Working paper*.
- [22] Sklar, A. (1959). Fonctions de repartition a n dimensions et leurs marges. *Publications de l'Institut de Statistique de Paris*, 8, 229–231.
- [23] Strasman, S., Ocello, A., Boyer, C., Le Corff, S., & Lemaire, V. (2024). An analysis of the noise schedule for score-based generative models. *Transactions on Machine Learning Research (TMLR)*.
- [24] Thinking Ahead Institute. (2025). *TPA Report : From Vision to Execution*. WTW / CAIA.
- [25] Wiese, M., Knobloch, R., Korn, R., & Kretschmer, P. (2020). Quant GANs : Deep generation of financial time series. *Quantitative Finance*, 20(9), 1419–1440.
- [26] Yoon, J., Jarrett, D., & van der Schaar, M. (2019). Time-series generative adversarial networks. *Advances in Neural Information Processing Systems*, 32.